

# Excercise\_6\_solutions\_part\_2

May 30, 2019

## 1 PART 2 - Logistic Regression Project

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
ad_data = pd.read_csv('advertising.csv')
```

```
In [2]: ad_data.head()
```

```
Out[2]:
```

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	\
0	68.95	35	61833.90	256.09	
1	80.23	31	68441.85	193.77	
2	69.47	26	59785.94	236.50	
3	74.15	29	54806.18	245.89	
4	68.37	35	73889.99	225.58	

	Ad Topic Line	City	Male	Country	\
0	Cloned 5thgeneration orchestration	Wrightburgh	0	Tunisia	
1	Monitored national standardization	West Jodi	1	Nauru	
2	Organic bottom-line service-desk	Davidton	0	San Marino	
3	Triple-buffered reciprocal time-frame	West Terrifurt	1	Italy	
4	Robust logistical utilization	South Manuel	0	Iceland	

	Timestamp	Clicked on Ad
0	2016-03-27 00:53:11	0
1	2016-04-04 01:39:02	0
2	2016-03-13 20:35:42	0
3	2016-01-10 02:31:19	0
4	2016-06-03 03:36:18	0

```
In [3]: ad_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
```

```

Daily Time Spent on Site    1000 non-null float64
Age                        1000 non-null int64
Area Income                1000 non-null float64
Daily Internet Usage       1000 non-null float64
Ad Topic Line              1000 non-null object
City                      1000 non-null object
Male                      1000 non-null int64
Country                   1000 non-null object
Timestamp                  1000 non-null object
Clicked on Ad              1000 non-null int64
dtypes: float64(3), int64(3), object(4)
memory usage: 78.2+ KB

```

```
In [4]: ad_data.describe()
```

```

Out[4]:      Daily Time Spent on Site      Age      Area Income \
count      1000.000000  1000.000000  1000.000000
mean         65.000200   36.009000  55000.000080
std         15.853615    8.785562  13414.634022
min          32.600000   19.000000  13996.500000
25%          51.360000   29.000000  47031.802500
50%          68.215000   35.000000  57012.300000
75%          78.547500   42.000000  65470.635000
max          91.430000   61.000000  79484.800000

      Daily Internet Usage      Male      Clicked on Ad
count      1000.000000  1000.000000  1000.000000
mean         180.000100    0.481000    0.500000
std         43.902339    0.499889    0.500250
min         104.780000    0.000000    0.000000
25%         138.830000    0.000000    0.000000
50%         183.130000    0.000000    0.500000
75%         218.792500    1.000000    1.000000
max         269.960000    1.000000    1.000000

```


## 2 Exploratory Data Analysis

```

In [5]: sns.set_style('whitegrid')
        ad_data['Age'].hist(bins=30)
        plt.xlabel('Age')

```

```
Out[5]: Text(0.5, 0, 'Age')
```




Create a jointplot showing Area Income versus Age.

```
In [6]: sns.jointplot(x='Age',y='Area Income',data=ad_data)
```


```
/home/kamil/anaconda3/lib/python3.7/site-packages/scipy/stats/stats.py:1713: FutureWarning: Us
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

```
Out[6]: <seaborn.axisgrid.JointGrid at 0x7ff04c566f28>
```



Create a jointplot showing the kde distributions of Daily Time spent on site vs. Age.


```
In [7]: sns.jointplot(x='Age',y='Daily Time Spent on Site',data=ad_data,color='red',kind='kde',
```



Create a jointplot of 'Daily Time Spent on Site' vs. 'Daily Internet Usage'

```
In [8]: sns.jointplot(x='Daily Time Spent on Site',y='Daily Internet Usage',data=ad_data,color='red')
```

```
Out[8]: <seaborn.axisgrid.JointGrid at 0x7ff01b3bb198>
```

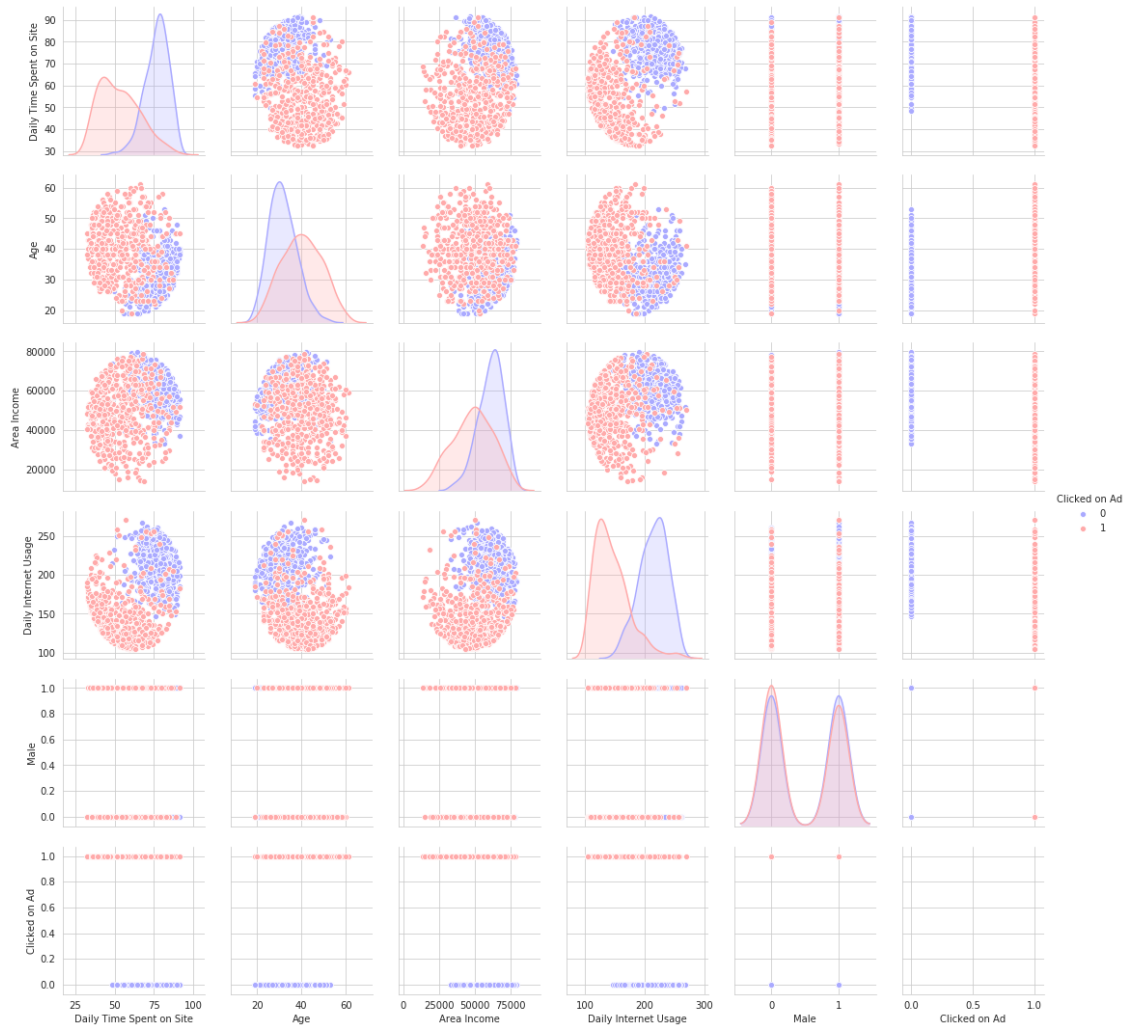


Finally, create a pairplot with the hue defined by the 'Clicked on Ad' column feature.

```
In [9]: sns.pairplot(ad_data, hue='Clicked on Ad', palette='bwr')
```

```
/home/kamil/anaconda3/lib/python3.7/site-packages/statsmodels/nonparametric/kde.py:488: RuntimeWarning:
  binned = fast_linbin(X, a, b, gridsize) / (delta * nobs)
/home/kamil/anaconda3/lib/python3.7/site-packages/statsmodels/nonparametric/kdetools.py:34: RuntimeWarning:
  FAC1 = 2*(np.pi*bw/RANGE)**2
/home/kamil/anaconda3/lib/python3.7/site-packages/numpy/core/fromnumeric.py:83: RuntimeWarning:
  return ufunc.reduce(obj, axis, dtype, out, **passkwargs)
```

```
Out[9]: <seaborn.axisgrid.PairGrid at 0x7ff01b2d1c88>
```



### 3 Logistic Regression

Now it's time to do a train test split, and train our model!

You'll have the freedom here to choose columns that you want to train on!

- Split the data into training set and testing set using `train_test_split`

```
In [10]: from sklearn.model_selection import train_test_split
```

```
In [11]: ad_data.head()
```

```
Out[11]:   Daily Time Spent on Site  Age  Area Income  Daily Internet Usage  \
0                68.95    35      61833.90                256.09
1                80.23    31      68441.85                193.77
2                69.47    26      59785.94                236.50
```

3	74.15	29	54806.18	245.89
4	68.37	35	73889.99	225.58

	Ad Topic Line	City	Male	Country \
0	Cloned 5thgeneration orchestration	Wrightburgh	0	Tunisia
1	Monitored national standardization	West Jodi	1	Nauru
2	Organic bottom-line service-desk	Davidton	0	San Marino
3	Triple-buffered reciprocal time-frame	West Terrifurt	1	Italy
4	Robust logistical utilization	South Manuel	0	Iceland

	Timestamp	Clicked on Ad
0	2016-03-27 00:53:11	0
1	2016-04-04 01:39:02	0
2	2016-03-13 20:35:42	0
3	2016-01-10 02:31:19	0
4	2016-06-03 03:36:18	0

```
In [12]: X = ad_data[['Daily Time Spent on Site', 'Age', 'Area Income', 'Daily Internet Usage',
y = ad_data['Clicked on Ad']
```

```
In [13]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=0)
```

Train and fit a logistic regression model on the training set.

```
In [14]: from sklearn.linear_model import LogisticRegression
```

```
In [15]: lm = LogisticRegression()
lm.fit(X_train, y_train)
```

```
/home/kamil/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning:
FutureWarning)
```

```
Out[15]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='warn',
n_jobs=None, penalty='l2', random_state=None, solver='warn',
tol=0.0001, verbose=0, warm_start=False)
```

## 4 Predictions and Evaluations

Now predict values for the testing data.

```
In [23]: y_pred = lm.predict(X_test)
```

Create a classification report for the model.

```
In [24]: from sklearn.metrics import classification_report
```

```
In [25]: print(classification_report(y_test, y_pred))
```



	precision	recall	f1-score	support
0	0.84	0.97	0.90	146
1	0.96	0.82	0.89	154
micro avg	0.89	0.89	0.89	300
macro avg	0.90	0.90	0.89	300
weighted avg	0.90	0.89	0.89	300

In [ ]: